

AI-POWERED FAKE NEWS DETECTION

A Comparative Study of Parametric vs. Non-Parametric Machine Learning Models

PROJECT 1 — FINAL DOCUMENTATION REPORT

2026

Author / Student	Vishnu Thangaraj
Enrollment ID / Roll No.	845488
Domain / Area	Natural Language Processing & Machine Learning
Project Period	June 2026 – July 2026 (6 Weeks)
Host Institution	IICT Summer Internship Programme 2026

EXECUTIVE SUMMARY

ABSTRACT

The proliferation of online news outlets and social media platforms has drastically increased the speed and volume of information dissemination while creating significant vulnerabilities to misinformation and deliberate disinformation. Automated fake news detection using machine learning provides a critical scalable countermeasure. In this study, an end-to-end natural language processing (NLP) classification pipeline was developed and evaluated on a benchmark dataset of 1,200 news articles collected across six international outlets (BBC, Reuters, CNN, Al Jazeera, The Guardian, The New York Times) and six thematic domains.

After text preprocessing and Term Frequency-Inverse Document Frequency (TF-IDF) feature extraction (capped at 5,000 maximum features), four prominent classification algorithms — K-Nearest Neighbours (KNN), Logistic Regression, Random Forest, and a Multi-Layer Perceptron (MLP) Neural Network — were trained on 960 articles and evaluated on an independent test set of 240 articles.

Parametric models (Logistic Regression and MLP) achieved a perfect 100.00% accuracy (F1-score 1.00), demonstrating high linear separability in the feature space. Non-parametric models also demonstrated robust performance (KNN: 97.50%, RF: 96.67%). This report presents a comprehensive comparative analysis of parametric versus non-parametric paradigms, evaluates predictive feature importance, details deployment considerations, and outlines future contextual embedding extensions for advanced fake news detection.

Keywords: Fake news detection, text classification, TF-IDF vectorisation, K-Nearest Neighbours, Logistic Regression, Random Forest, Multi-Layer Perceptron, natural language processing, machine learning, misinformation, disinformation, text mining, artificial intelligence.

TABLE OF CONTENTS

1	Introduction (Background)	Page 4
2	Introduction (Impact) & Problem Statement	Page 5
3	Literature Review (Part 1)	Page 6
4	Literature Review (Part 2 & Comparison)	Page 7
5	Proposed Methodology	Page 8
6	Dataset Description (Source & Overview)	Page 9
7	Dataset Description (Schema & Distribution)	Page 10
8	Data Preprocessing (Concepts)	Page 11
9	Data Preprocessing (Implementation)	Page 12
10	Feature Engineering (TF-IDF Theory)	Page 13
11	Feature Engineering (Implementation)	Page 14
12	Machine Learning Models (Linear & Probabilistic)	Page 15
13	Machine Learning Models (Ensembles & NNs)	Page 16

1. INTRODUCTION

1.1 Background Context

In the digital age, news consumption has dramatically shifted from traditional print and broadcast media to online social networks, news aggregators, and digital messaging platforms. While this technological transition has democratised information access globally, significantly lowering the barrier to entry for content publication, it has simultaneously created a fertile ecosystem for the rapid spread of false, misleading, and manipulated content — colloquially and broadly known as “fake news.”

Fake news encompasses a wide spectrum of content. At the benign end of the spectrum is satire and parody, intended to humorously critique society without malicious intent. However, at the malignant end lies deliberate disinformation: outright fabrications explicitly designed to deceive readers for specific political, financial, or social gain. This includes state-sponsored propaganda, financially motivated clickbait generated by ad-revenue farms, and ideologically driven conspiracy theories.

The manual identification of fake news is labour-intensive, highly subjective, and fundamentally unscalable. Human fact-checkers, even when highly trained and well-resourced, cannot keep pace with the sheer volume of content generated every second across platforms like X (formerly Twitter), Facebook, and Reddit. Consequently, there is an urgent and critical need for automated solutions.

Artificial Intelligence (AI), specifically Natural Language Processing (NLP) and Machine Learning (ML), provides the mathematical framework necessary to analyze linguistic patterns at a massive scale. By treating text as data, algorithms can identify structural anomalies, emotional manipulation tactics, and source discrepancies much faster than human cognition allows. This study leverages these AI methodologies to build a robust fake news detection pipeline.

1.2 Impact of Fake News

The societal consequences of fake news are far-reaching, crossing international borders and impacting various critical domains of human life:

- **Public Health:** During the COVID-19 pandemic, health misinformation caused measurable harm globally. The rapid spread of false cures and anti-vaccination conspiracies directly led to increased mortality, dangerous self-medication practices, and systemic non-compliance with public safety guidelines.
- **Democratic Integrity:** False narratives regarding electoral processes, candidate histories, and voter fraud undermine public trust in democratic institutions. Disinformation campaigns are frequently utilized to suppress voter turnout and polarize the electorate.
- **Economic Stability:** Market-moving misinformation about corporations, financial institutions, or national economies can cause severe stock volatility, triggering automated trading algorithms and causing investor panic, resulting in massive capital loss.

2. PROBLEM STATEMENT & OBJECTIVES

2.1 Problem Statement

The primary problem this project addresses is binary text classification — computationally labelling each news article as either *real* (label = 0) or *fake* (label = 1). The massive scale of unstructured textual data necessitates a supervised machine learning approach capable of parsing latent semantic, lexical, and syntactical patterns to distinguish credible journalism from fabricated text.

2.2 Project Objectives

1. To develop a complete, end-to-end text classification pipeline in Python utilizing `scikit-learn` for automated fake news detection.
2. To statistically compare the efficacy of parametric models (Logistic Regression, Multi-Layer Perceptron, Naive Bayes) against non-parametric ensemble models (Random Forest, K-Nearest Neighbours) using a standardized benchmark.
3. To rigorously evaluate these models using established NLP performance metrics: accuracy, precision, recall, F1-score, and confusion matrices.
4. To analyze model interpretability, determine predictive feature importance, and construct a highly optimized inference module ready for deployment.

3. LITERATURE REVIEW

3.1 Evolution of Fake News Detection

Fake news detection has been extensively studied over the past decade, evolving rapidly from simple keyword matching algorithms to highly complex, multi-modal deep learning architectures. Academic research has largely bifurcated into two primary approaches: content-based analysis (focusing strictly on the text and media of the article) and context-based analysis (focusing on how the article spreads through a social network and user engagement).

In the realm of content-based analysis, **Conroy et al. (2015)** proposed foundational linguistic and network-based methods. Their work highlighted the importance of lexical features, syntax, and sentiment markers. Their research indicated that fake news often diverges stylistically from traditional journalistic standards, frequently employing hyperbolic language and emotional appeals. This established the viability of using Natural Language Processing (NLP) to extract semantic meaning from raw text.

Perez-Rosas et al. (2018) expanded on this by developing a novel dataset focused on celebrity and political news. They demonstrated the profound effectiveness of linear Support Vector Machines (SVMs) combined with advanced linguistic feature engineering. Their findings suggested that when text features are meticulously crafted (e.g., n-grams, POS tagging), simple linear models can achieve state-of-the-art results, often outperforming complex, un-tuned neural networks.

Conversely, focusing on context-based analysis, **Shu et al. (2017)** provided a comprehensive survey emphasizing the role of user engagements and social context in verification. They argued that because fake news is designed to go viral, the meta-data surrounding an article's propagation (retweets, user comments, network topology) is just as predictive, if not more so, than the text itself.

3.2 Deep Learning and Hybrid Approaches

As computational power increased, research shifted heavily toward deep learning. **Wang (2017)** introduced the LIAR dataset, a highly popular benchmark composed of short statements evaluated by PolitiFact. Wang evaluated various deep learning models, including Convolutional Neural Networks (CNNs) and Bidirectional Long Short-Term Memory networks (BiLSTMs). The study proved that deep learning architectures are superior at capturing the sequential and contextual nuances of text compared to traditional bag-of-words approaches.

To capture both text and social propagation, **Ruchansky et al. (2017)** proposed the CSI (Capture, Score, and Integrate) model. This hybrid architecture integrates the textual representation of the article, the response of users on social media, and the historical behaviour of the source outlet. This tri-partite approach yielded highly accurate detections, validating that multi-modal signals are the future of robust verification.

More recently, **Kaliyar et al. (2020)** developed FNDNet, a deep convolutional neural network specifically tailored for fake news. Their extensive testing on the BuzzFeed and PolitiFact datasets proved that deep automated feature extraction generally outperforms manual feature engineering in massive, unconstrained datasets, though at the cost of high computational overhead and reduced interpretability.

3.3 Literature Comparison Table

Reference / Paper	Methodology Focus	Dataset Utilized	Key Findings
Conroy et al. (2015)	Linguistic & Syntax	Custom / Crowdsourced	Lexical anomalies strongly indicate fake news.
Shu et al. (2017)	Social Networks	Twitter / Facebook API	Engagement metrics are crucial for early detection.
Perez-Rosas et al. (2018)	Linear SVMs	Celebrity / Political	Linear models excel on well-engineered text features.
Wang (2017)	BiLSTM / CNN	LIAR Benchmark Dataset	Deep networks successfully capture sequential context.
Ruchansky et al. (2017)	CSI Hybrid Model	Weibo / Twitter	Combining user behaviour with text improves accuracy.
Kaliyar et al. (2020)	FNDNet (CNN)	BuzzFeed / PolitiFact	Deep feature extraction outperforms manual tuning.

4. PROPOSED METHODOLOGY

The proposed methodology relies on a highly structured, sequential workflow designed to process raw textual strings, extract mathematical representations of language, and train discriminative classification models. This modular pipeline ensures seamless integration into production environments, allowing for hot-swapping of models or vectorisers without breaking the underlying architecture.

4.1 Complete Workflow Architecture



5. DATASET DESCRIPTION

5.1 Data Source and Collection

The dataset utilized in this project is a robust, curated benchmark specifically constructed for the IICT Summer Internship Programme 2026. It aggregates news article metadata and raw body text scraped and compiled from six prominent international news outlets, representing a diversity of journalistic styles and regional biases:

- **BBC News** (United Kingdom)
- **Reuters** (Global)
- **CNN** (United States)
- **Al Jazeera** (Middle East / Global)
- **The Guardian** (United Kingdom)
- **The New York Times** (United States)

To ensure the classification algorithms do not overfit to a single domain of discourse (e.g., politics), the dataset strategically covers six distinct thematic categories: Politics, Health, Business, Entertainment, Sports, and Technology. This guarantees that the fake news models learn stylistic linguistic features of deception rather than simply memorizing the vocabulary of a single topic.

5.2 Dataset Size and Integrity

The total dataset comprises exactly **1,200 unique records** spanning **14 distinct feature columns**. A comprehensive data integrity check was executed upon loading the dataset into pandas.

Thorough data validation confirmed that no missing values (nulls or NaNs) were present in the primary classification fields (`full_text` for the input variables, and `label` for the target variable). This high level of dataset cleanliness ensured complete data retention for modelling without requiring complex statistical imputation techniques or row deletion.

5.3 Dataset Schema and Columns


While the `full_text` column serves as the primary driver for the NLP pipeline, the dataset provides a rich array of metadata that can be utilized in future hybrid models.

Column Name	Data Type	Description
<code>id</code>	Integer	Unique primary key identifier for the article
<code>title</code>	String	Headline text of the article
<code>full_text</code>	String	Primary NLP feature — the complete article body text
<code>source_name</code>	Categorical	Publisher outlet (e.g., BBC, Reuters)
<code>publish_date</code>	Datetime	The date the article was published online
<code>category</code>	Categorical	Politics, Health, Business, Entertainment, Sports, Tech
<code>label</code>	Integer (0/1)	Target variable: 0 = Real News, 1 = Fake News
<code>sentiment_score</code>	Float	Calculated sentiment polarity score [-1.0, 1.0]
<code>clickbait_score</code>	Float	Probability metric indicating headline sensationalism
<code>readability_score</code>	Float	Flesch Reading Ease score evaluating text complexity

Table II. Comprehensive Dataset Schema

5.4 Class Distribution and Balance

Exploratory Data Analysis (EDA) was conducted to inspect the target label distribution. This is a vital step because severe class imbalance (e.g., 90% real, 10% fake) mathematically skews standard accuracy metrics and forces models to trivially predict the majority class.



6. DATA PREPROCESSING

6.1 Preprocessing Concepts

Data preprocessing is the foundational bedrock of any Natural Language Processing pipeline. Raw text data pulled from the internet is inherently noisy, unstructured, and filled with grammatical artifacts that offer absolutely no predictive value but drastically increase the computational dimensionality of the feature space.

The preprocessing strategy implemented in this project involves multiple sequential standardisation operations, carefully designed to distill the text down to its most meaningful semantic roots:

- **Lowercase Conversion:** All text is strictly converted to lowercase. Algorithms treat strings with different cases as entirely separate tokens. Lowercasing ensures absolute uniformity (e.g., 'Government', 'government', and 'GOVERNMENT' map to the identical token).
- **HTML & URL Removal:** HTML markup tags (`<br>` , `<a>`) and hyperlinks (HTTP/HTTPS) are stripped using targeted regular expressions. URL metadata does not contribute to the linguistic narrative and creates highly sparse, unique tokens.
- **Punctuation & Special Characters:** Punctuation marks, numbers, and special characters (e.g., @, #, \$, %) are purged to isolate purely lexical alphabetic words.
- **Stopwords Removal:** Extremely common English stopwords (e.g., 'the', 'is', 'in', 'and', 'of') are filtered out using NLTK's standard stopwords corpus. While syntactically necessary for human readers, they carry negligible discriminative power for classification and simply bloat the matrix.
- **Tokenization & Lemmatization:** The continuous string of text is tokenized into a list of discrete word units. These units are then lemmatized to their base dictionary form. For example, 'running', 'ran', and 'runs' are conflated to the root lemma 'run'. This step drastically reduces vocabulary sparsity.

6.2 Preprocessing Implementation (Code Snippet)

The conceptual preprocessing steps were implemented in Python utilizing the built-in `re` (regex) module alongside functions provided by the Natural Language Toolkit (NLTK).

```
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer

# Download required NLTK resources
nltk.download('stopwords')
nltk.download('wordnet')

# Initialize lemmatizer and stop words list
lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    """
    Cleans raw text for NLP classification by removing noise,
    stopwords, and lemmatizing the remaining tokens.
    """
    # 1. Type coercion and lowercase
    text = str(text).lower()

    # 2. Remove URLs
    text = re.sub(r'https?://\S+|www\.\S+', '', text)

    # 3. Remove HTML tags
    text = re.sub(r'<.*?>', '', text)

    # 4. Remove punctuation and numbers (keep only alphabets)
    text = re.sub(r'[^a-zA-Z\s]', '', text)

    # 5. Tokenization, Stopwords removal, and Lemmatization
    tokens = text.split()
    cleaned_tokens = [
        lemmatizer.lemmatize(word)
        for word in tokens
        if word not in stop_words
    ]
```

By applying this function, a noisy sentence like `"BREAKING NEWS!!! The government is planning a massive coverup in 2026. Click here: http://fake.url"` is seamlessly transformed into a clean, model-ready string: `"breaking news government planning massive coverup"`.

7. FEATURE ENGINEERING

7.1 The Bag-of-Words Paradigm

Machine learning models cannot natively process raw string text; they require numerical input. Therefore, following preprocessing, feature engineering is required to transform the clean text corpus into a structured numerical matrix. This is commonly referred to as text vectorisation.

While simplistic approaches like Count Vectorisation (Boolean Bag-of-Words) exist, they merely count the raw frequency of words, ignoring the relative importance of a word to the specific document compared to the global corpus. In this project, the **Term Frequency-Inverse Document Frequency (TF-IDF)** method is employed exclusively to provide a more sophisticated statistical weighting.

7.2 TF-IDF Mathematical Framework

TF-IDF evaluates the relevance of a word to a document within a broader corpus. It is calculated by multiplying two distinct metrics:

1. Term Frequency (TF): Measures how frequently a term (t) occurs in a document (d). Because documents have different lengths, it is often normalized by the total number of words in the document.

$$TF(t, d) = (\text{Count of } t \text{ in } d) / (\text{Total words in } d)$$

2. Inverse Document Frequency (IDF): Measures how important a term is globally. It weighs down highly frequent terms across the corpus and scales up rare terms.

$$IDF(t, D) = \log_e (\text{Total documents} / \text{Documents containing } t)$$

The final TF-IDF score is the product of these two metrics. Words that are ubiquitously common across all documents are severely penalized (score nears zero), while words that are highly frequent in a specific document but rare globally are assigned heavy mathematical weights, signifying their discriminative importance for that specific text.

7.3 Dimensionality Control & Implementation

A major challenge with TF-IDF is the "curse of dimensionality." If a corpus contains 100,000 unique words, the resulting matrix will have 100,000 columns, causing memory bloat and severe model overfitting. To control computational complexity, the maximum number of features (vocabulary size) is capped at exactly **5,000** using the `max_features` parameter in scikit-learn.

```
from sklearn.feature_extraction.text import TfidfVectorizer

# Initialize the TF-IDF Vectoriser
# We restrict the vocabulary to the top 5,000 most frequent terms
# We also include bigrams (ngram_range=(1,2)) to capture 2-word phrases
tfidf_vectorizer = TfidfVectorizer(
    max_features=5000,
    ngram_range=(1, 2),
    stop_words='english'
)

# Note: The vectoriser is FITTED ONLY ON THE TRAINING SET
# to strictly prevent data leakage from the test set.
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)

# The testing set is TRANSFORMED using the already fitted vocabulary
X_test_tfidf = tfidf_vectorizer.transform(X_test)

print(f"Train Matrix Shape: {X_train_tfidf.shape}")
print(f"Test Matrix Shape: {X_test_tfidf.shape}")
# Output:
# Train Matrix Shape: (960, 5000)
# Test Matrix Shape: (240, 5000)
```

The resulting `X_train_tfidf` is a highly sparse SciPy matrix. Sparse matrices optimize memory by only storing non-zero values, which is critical since the vast majority of documents will only contain a tiny fraction of the 5,000 words in the global vocabulary.

8. MACHINE LEARNING MODELS

This study evaluates four diverse classification algorithms representing fundamental machine learning paradigms: parametric linear models, probabilistic models, non-parametric ensemble models, and parametric neural networks. Evaluating multiple paradigms ensures a comprehensive analysis of how different mathematical approaches handle high-dimensional TF-IDF data.

8.1 Logistic Regression (Linear Parametric)

Working Mechanism: Despite its name, Logistic Regression is a classification algorithm. It estimates the probability of a binary outcome (0 or 1) using a logistic sigmoid function. Mathematically, it attempts to fit a linear decision boundary (a multi-dimensional hyperplane) across the 5,000-dimensional feature space to perfectly separate the "Real" documents from the "Fake" documents.

Advantages: Logistic Regression offers exceptional computational efficiency, requiring minimal memory and training in milliseconds. Crucially, it boasts high interpretability; developers can directly extract the coefficients assigned to each word to understand exactly how the model makes its decisions.

Disadvantages: It possesses a fundamental inability to capture complex, non-linear relationships without manual, explicit feature interaction engineering.

8.2 Naive Bayes (Probabilistic)

Working Mechanism: A probabilistic classifier grounded in Bayes' Theorem. It relies on the "naive" assumption of conditional independence between every pair of features (words) given the class label. For TF-IDF feature matrices containing fractional weights, the Multinomial variant (MultinomialNB) is mathematically suitable.

Advantages: It provides extremely fast training and inference speeds, is highly scalable to massive datasets, and traditionally functions as a remarkably strong baseline for document classification tasks even with limited training data.

Disadvantages: The core assumption — that all words appear completely independently of one another — is inherently violated in human language syntax, sometimes leading to overconfident but incorrect probability estimations.

8.3 Random Forest (Non-Parametric Ensemble)

Working Mechanism: Random Forest is a non-parametric ensemble learning method. During training, it constructs a multitude of decision trees (an ensemble) using a technique called bootstrap aggregating (bagging). Crucially, at each split in the tree, it only considers a random subset of the 5,000 features. The final prediction is determined by taking the mode (majority vote) of the individual tree predictions.

Advantages: Random Forest naturally handles high-dimensional data exceptionally well. Because it averages many independent trees, it is highly resilient to overfitting on training data and is capable of modeling highly complex non-linear decision boundaries.

Disadvantages: It requires a substantially higher memory footprint to store the ensemble of trees and suffers from longer inference times. Furthermore, it acts largely as a "black box" compared to linear models, making it much harder to extract clear lexical rules.

8.4 Multi-Layer Perceptron (Neural Network)

Working Mechanism: An MLP is a parametric, feed-forward artificial neural network. The architecture utilized in this study consists of an input layer with 5,000 nodes (one for each TF-IDF feature), a single dense hidden layer containing 100 neurons equipped with ReLU activation functions, and a single binary output neuron using a sigmoid activation. It relies on backpropagation and gradient descent to optimize internal weights.

Advantages: Neural networks possess exceptional predictive capacity. The Universal Approximation Theorem dictates that they have the theoretical capacity to approximate any continuous non-linear function, allowing them to capture deeply hidden hierarchical patterns in the text.

Disadvantages: Severe lack of interpretability, high sensitivity to hyperparameter selection (e.g., learning rate, architecture depth), and computationally intensive training requirements that often demand GPU acceleration for larger datasets.

9. MODEL TRAINING

Model training, hyperparameter configuration, and validation were conducted uniformly using the `scikit-learn` (sklearn) library in Python. A standardized, rigorous approach was maintained across all algorithms to ensure a fair and scientifically valid comparative analysis.

9.1 Train-Test Split (80:20 Stratified)

The entire dataset of 1,200 records was partitioned into training and testing subsets using an 80/20 ratio. This split was executed using a **stratified** sampling technique. Stratification ensures that the exact 50/50 balance of Real and Fake news articles present in the global dataset is proportionally maintained in both the training and testing subsets.

Furthermore, a fixed random seed (`random_state=42`) was hardcoded during the split. This guarantees complete reproducibility of the shuffle and split across multiple experimental iterations by any secondary researcher. The split resulted in a training set of 960 articles and an independent, completely held-out testing set of 240 articles.

9.2 Library Selection and Hyperparameters

The models were imported from `scikit-learn` and instantiated with specific, optimized configurations:

- **Logistic Regression:** Employed the L-BFGS (Limited-memory Broyden-Fletcher-Goldfarb-Shanno) optimization solver, with a maximum of 1,000 iterations to ensure full gradient convergence without triggering timeout warnings.
- **Naive Bayes:** Utilized the `MultinomialNB` class with default Laplace smoothing (`alpha=1.0`) to handle any out-of-vocabulary features gracefully.
- **Random Forest:** Initialized with an ensemble size of 100 decision estimators (`n_estimators=100`), utilizing Gini impurity for split quality evaluation, and constrained by `random_state=42`.
- **Neural Network:** Configured with a single hidden layer composed of 100 neurons (`hidden_layer_sizes=(100,)`), utilizing the Adam optimizer, early stopping disabled, and a maximum of 300 epochs.

9.3 Model Training Implementation (Code Snippet)

The following Python code demonstrates the exact implementation used to instantiate, train, and generate predictions for all four evaluated algorithms over the TF-IDF feature matrices.

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score

# 1. Stratified Train-Test Split (80% Train, 20% Test)
X_train, X_test, y_train, y_test = train_test_split(
    df['clean_text'],
    df['label'],
    test_size=0.2,
    stratify=df['label'],
    random_state=42
)

# (TF-IDF Vectorisation occurs here, resulting in X_train_tfidf,
X_test_tfidf)

# 2. Dictionary of instantiated models with hyperparameters
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000,
solver='lbfgs'),
    "Naive Bayes": MultinomialNB(alpha=1.0),
    "Random Forest": RandomForestClassifier(n_estimators=100,
random_state=42),
    "Neural Network (MLP)": MLPClassifier(hidden_layer_sizes=(100,),
max_iter=300)
}

# 3. Training and Inference Loop
predictions_dict = {}

for name, model in models.items():
    print(f"Training {name}...")

    # Fit the model on the 5,000-feature sparse training matrix
    model.fit(X_train_tfidf, y_train)

    # Generate predictions on the unseen testing matrix
```

This loop structure ensures that all algorithms are trained and evaluated on the exact same data splits, providing a strictly controlled environment for comparative analysis.

10. RESULTS & EVALUATION

The fully trained models were rigorously evaluated against the independent, unseen testing set of 240 articles. Standard binary classification metrics were calculated using `sklearn.metrics.classification_report`.

The parametric models (Logistic Regression and MLP) achieved flawless predictive performance, indicating that the text dataset is highly linearly separable when mapped into the 5,000-dimensional TF-IDF vector space.

10.1 Evaluation Metrics Table

Algorithm	Accuracy	Precision	Recall	F1-Score
Logistic Regression	100.00%	1.00	1.00	1.00
Neural Network (MLP)	100.00%	1.00	1.00	1.00
Naive Bayes	98.33%	0.99	0.98	0.98
Random Forest	96.67%	0.97	0.97	0.97
K-Nearest Neighbours*	97.50%	0.98	0.97	0.97

Table III. Model Performance Comparison (Test Set n=240)

* KNN provided for extended baseline comparison.

10.2 Metric Definitions

- **Accuracy:** The total percentage of correct predictions out of all test samples. Because the test set is perfectly balanced, accuracy is a highly reliable core metric.
- **Precision (Fake News):** Of all the articles the model flagged as Fake, what percentage were actually Fake? High precision means minimizing "False Alarms" (flagging real news as fake).
- **Recall (Fake News):** Of all the actual Fake news articles in the dataset, what percentage did the model successfully find? High recall means minimizing "Missed Fakes".
- **F1-Score:** The harmonic mean of Precision and Recall, providing a balanced, single-value metric that punishes extreme disparities between the two.

10.3 Confusion Matrices & Comparison Graph

Confusion matrices visually map the exact distribution of True Positives, True Negatives, False Positives (Type I errors), and False Negatives (Type II errors). The matrices validate that Logistic Regression exhibited perfect classification (0 errors), whereas Random Forest yielded 8 total misclassifications across the 240 test samples.




Fig. 2 - Logistic Regression CM (0 Errors)




Fig. 3 - Random Forest CM (8 Errors)




Fig. 4 - Visual Accuracy Comparison Chart across Evaluated Models

The flawless performance of linear models (LR) suggests that complex, non-linear algorithms (like Random Forest) may actually be slightly overfitting to noise in the training set, slightly degrading their generalisation to the test set compared to a simple, mathematically smooth hyperplane.

11. FEATURE IMPORTANCE & ANALYSIS

11.1 Extracting Lexical Markers

Model interpretability is paramount in fake news detection to understand *why* an algorithm made a specific classification. By inspecting the internal mathematical weights (coefficients) of the Logistic Regression model, we can rank the precise words from the 5,000-term TF-IDF vocabulary that most strongly pull a prediction toward "Fake" (positive coefficients) or "Real" (negative coefficients).

```
import pandas as pd

# Extract the fitted feature names from the TF-IDF Vectoriser
feature_names = tfidf_vectorizer.get_feature_names_out()

# Extract the learned coefficients from the Logistic Regression model
coefficients = models["Logistic Regression"].coef_[0]

# Create a DataFrame mapping words to their importance weights
feature_importance = pd.DataFrame({
    'Word': feature_names,
    'Weight': coefficients
})

# Sort to find the strongest Fake News indicators
fake_indicators = feature_importance.sort_values(by='Weight',
ascending=False).head(10)

# Sort to find the strongest Real News indicators
real_indicators = feature_importance.sort_values(by='Weight',
ascending=True).head(10)
```

11.2 Lexical Analysis Results

Fake News Indicators: Words heavily associated with fake news typically include sensationalist, highly emotive, or hyperbolic terms. Examples extracted from this dataset include *'hoax'*, *'shocking'*, *'conspiracy'*, *'breaking'*, and highly polarized political rhetoric used without factual context. These linguistic markers attempt to elicit a strong emotional response rather than convey objective facts.

Real News Indicators: Conversely, words mathematically associated with real news tend to be neutral, objective, and specific to verified entities, dates, locations, and journalistic sources. Examples include *'reuters'*, *'announced'*, *'wednesday'*, *'spokesperson'*, and *'statement'*.

12. PREDICTION MODULE

12.1 Deployable Inference Function

The final operational phase of the project encapsulates the best-performing model (Logistic Regression) alongside the fitted TF-IDF vectoriser into a deployable Python inference pipeline. This module accepts raw, unseen string text from a user, executes the exact preprocessing regimen, vectorises the text, and outputs a classification accompanied by a confidence probability.

```
def predict_fake_news(raw_text):  
    '''  
    End-to-end inference function for deployment.  
    '''  
    # 1. Apply the exact same preprocessing used in training  
    clean_text = preprocess_text(raw_text)  
  
    # 2. Vectorise using the pre-fitted TF-IDF model  
    vectorized_text = tfidf_vectorizer.transform([clean_text])  
  
    # 3. Generate prediction and confidence probability  
    model = models["Logistic Regression"]  
    prediction = model.predict(vectorized_text)[0]  
    probabilities = model.predict_proba(vectorized_text)[0]  
  
    label = "FAKE NEWS" if prediction == 1 else "REAL NEWS"  
    confidence = probabilities[prediction] * 100
```

12.2 Sample Input & Output Demonstration

[USER INPUT]: "BREAKING: Alien spacecraft found in Area 51, government confirms massive coverup!"

[OUTPUT PREDICTION]: FAKE NEWS

Confidence: 99.82%

[USER INPUT]: "The Federal Reserve announced a new interest rate hike today to combat rising inflation metrics."

[OUTPUT PREDICTION]: REAL NEWS

Confidence: 98.45%

13. ADVANTAGES, LIMITATIONS & FUTURE SCOPE

13.1 Advantages of the Proposed System

The developed machine learning pipeline offers several significant advantages for immediate production deployment. Foremost, it is highly accurate and computationally efficient. Models like Logistic Regression require a very small memory footprint and offer inference times measured in single-digit milliseconds, making them ideal for high-throughput REST APIs, mobile backends, or real-time browser extensions. Furthermore, the model is highly interpretable; developers can easily audit the feature coefficients to understand exactly which words trigger a classification, ensuring system transparency and trust.

13.2 Technical Limitations

Despite its high accuracy on the benchmark, the system possesses notable architectural limitations. It relies entirely on a "bag-of-words" (TF-IDF) representation, which fundamentally ignores word order, grammatical structure, and semantic context (e.g., the model cannot distinguish between the sentences "dog bites man" and "man bites dog"). The flawless 100% accuracy on the test set suggests that the dataset may contain stylistic artifacts that make the classification artificially straightforward for linear algebra. This raises valid concerns about the model's ability to generalize to highly sophisticated, adversarial disinformation campaigns written by advanced Generative AI (like ChatGPT), which are explicitly designed to mimic the neutral, objective tone of real journalism.

13.3 Future Scope & Enhancements

Future scope for this project involves migrating away from pure TF-IDF to advanced, context-aware deep learning embeddings, such as BERT (Bidirectional Encoder Representations from Transformers) or RoBERTa. These Transformer architectures intrinsically understand semantic context, sarcasm, and sentence-level nuance.

Additionally, modern disinformation is rarely limited to pure text. Integrating multi-modal detection signals — such as image forensics to detect manipulated photos or deepfakes, source URL domain credibility scoring, and social network propagation graphs to analyze how rapidly an article is shared — would provide a significantly more robust, defense-in-depth approach against evolving, complex fake news operations.

14. CONCLUSION

In conclusion, this project successfully developed, trained, and rigorously evaluated an end-to-end supervised machine learning pipeline for the automated detection of fake news. Utilizing a diverse benchmark dataset of 1,200 news articles spanning multiple geographic regions and topical categories, the study applied strict textual preprocessing combined with Term Frequency-Inverse Document Frequency (TF-IDF) feature extraction to convert unstructured human language into structured numerical data suitable for algorithmic ingestion.

The comprehensive comparative analysis between parametric algorithms (Logistic Regression, Multi-Layer Perceptron) and non-parametric algorithms (Random Forest) yielded decisive, empirical results. The parametric linear models achieved perfect, flawless accuracy (100.00%) on the unseen testing set, demonstrating the incredibly high linear separability of the dataset when projected into the 5,000-dimensional TF-IDF feature space. While complex non-parametric models performed admirably, they were marginally outperformed in this specific bag-of-words paradigm, likely due to slight overfitting.

Furthermore, an in-depth analysis of feature importance confirmed that fake news articles frequently rely on distinct linguistic markers — highly emotive, sensationalist, and hyperbolic vocabulary — which algorithms can reliably detect and quantify. The project emphatically highlights the efficacy of fundamental NLP techniques in identifying and combatting misinformation at scale.

While the current bag-of-words approach proved highly effective, interpretable, and computationally optimal for this dataset, continuous algorithmic iteration is required. As malicious actors increasingly utilize generative Artificial Intelligence to craft sophisticated disinformation, future detection pipelines must evolve to leverage deep contextual transformers. Ultimately, this AI-powered detection system serves as a highly robust, scalable foundational tool in the ongoing global effort to preserve information integrity, public trust, and democratic stability in the modern digital age.

15. REFERENCES

- [1] N. J. Conroy, V. L. Rubin, and Y. Chen, "Automatic deception detection: Methods for finding fake news," *Proc. Assoc. Info. Science and Tech.*, vol. 52, pp. 1-4, 2015.
- [2] K. Shu, et al., "Fake news detection on social media: A data mining perspective," *ACM SIGKDD Explorations*, vol. 19, pp. 22-36, 2017.
- [3] V. Perez-Rosas, et al., "Automatic detection of fake news," *arXiv preprint arXiv:1708.07104*, 2017.
- [4] W. Y. Wang, "'Liar, liar pants on fire': A new benchmark dataset for fake news detection," *arXiv:1705.00648*, 2017.
- [5] N. Ruchansky, S. Seo, and Y. Liu, "Csi: A hybrid deep model for fake news detection," *CIKM 2017*, pp. 797-806, 2017.
- [6] R. K. Kaliyar, et al., "FNDNet-a deep convolutional neural network for fake news detection," *Cognitive Systems Research*, vol. 61, pp. 32-44, 2020.

APPENDIX A. ENVIRONMENT DEPENDENCIES

```
# required_packages.txt
pandas==2.2.0
numpy==1.26.4
scikit-learn==1.4.0
nltk==3.8.1
matplotlib==3.8.2
```

APPENDIX B. GLOSSARY OF TERMS

- **NLP:** Natural Language Processing.
- **TF-IDF:** Term Frequency - Inverse Document Frequency.
- **Lemmatization:** Grouping together the inflected forms of a word.
- **Parametric Model:** A learning model that summarizes data with a set of parameters of fixed size (e.g. Logistic Regression).
- **Non-Parametric Model:** A model where the number of parameters grows with the size of the training data (e.g. KNN, Random Forest).